

**GEORGE HALAL** | Personal Site: <https://georgehalal.github.io> | [georgehalal@alumni.stanford.edu](mailto:georgehalal@alumni.stanford.edu) | +1 (650) 422-9033

Stanford PhD turned AI engineer with expertise in agentic search, synthetic data generation, and LLM finetuning.  
Most recently, I trained and open-sourced rerankers on the cost/performance frontier across public and customer benchmarks.

EDUCATION

Stanford University	Ph.D. Physics	GPA: 4.00/4.00	June 2019–July 2024
Lehigh University	B.S. Physics & Minor in Applied Mathematics	GPA: 3.97/4.00	Aug. 2015–May 2019

EXPERIENCE

**Member of Technical Staff** | Contextual AI, Mountain View, CA | July 2024—Present

- Best-performing, Most Efficient, Instruction-Following LLM Rerankers** | [v2 Blog](#) | [v1 Blog](#) | [Snowflake Blog](#) | [OS models](#) | [OS evals](#)
- Built a synthetic data pipeline to generate contrastive data covering desired behaviors and diverse domains.
  - Experimented with QAT, distillation, reinforcement learning, and curriculum learning among other techniques.
  - Achieved SOTA performance on instruction following, question answering, multilinguality, product search/recommendation systems, and real world use cases.
  - Selected as the default reranker for Snowflake Cortex AI among other companies.

- Agentic Search Tool Use Optimization**
- Optimized the type, number, cost, and latency of knowledge search and navigation tools an agent uses during rollouts.

- Graph-based Search (Graph RAG)** | Paper in Prep
- Developed an LLM-based pipeline to turn documents into knowledge graphs for efficient retrieval at query time.
  - Shipped to production as part of a mixture of retrievers for answering certain queries.
  - Separately, mentored a Stanford CS student on his master’s thesis, “End-to-End Retrieval on Black-Box Knowledge Graphs.”

**Graduate Student Researcher** | Stanford University, Stanford, CA | June 2019–July 2024

- Transformer-Based Super-Resolution for Dust Polarization Images** | [GitHub Link](#)
- Trained a multi-image encoder, a transformer-based fusion module, and a decoder to increase the image resolutions by 4x.
- Causal Inference for Modeling the Effects of the Nearby Dust Geometry on Magnetic Fields** | [Paper Link](#)
- Spherical Harmonic Convolutional Hough Transform** | [GitHub Link](#) | [Paper Link](#) | [Invited Talk Link](#)
- Achieved 3000x speedup and 5x memory reduction over the previous SOTA for modeling the structure of interstellar gas.
- Modeling the Foreground Obscuring Radiation from the Early Universe** | [Paper Link](#) | [Award Link](#) | Invited Talks: [Harvard](#), [Spain](#), [S4](#)
- Applied computer vision and Bayesian inference for quantifying this signal, setting new limits on early universe expansion.
- Conditional Wasserstein Generative Adversarial Network with Gradient Penalty for Generating Observed Galaxy Properties**
- Deep Learning for Modeling the Transfer Function of Galaxy Detection**, achieving an ROC-AUC score of 0.95 | [GitHub Link](#)
- Deep Learning for Searching for 2- $\nu$  Double- $\beta$  Decay of  $^{136}\text{Xe}$  to the Excited State of  $^{136}\text{Ba}$  in EXO-200 Data** | [Poster Link](#)
- Developed a data acquisition pipeline and an LSTM-based model, achieving an ROC-AUC score of 0.98.

**Data Scientist Intern** | Alife Health, San Francisco, CA | June 2023—Sept. 2023

Causal Inference, A/B Testing, and Machine Learning for IVF Intracycle Dose Adjustments

**Undergraduate Student Researcher** | Yale University and Lehigh University | Nov. 2016–May 2019

- Deep Learning for Heavy-Flavor Jet Classification at RHIC** | [Report Link](#) | [Talk Link](#)
- Deep Learning for Collision Geometry Determination**

SKILLS

Python • PyTorch • WandB • Pandas • vLLM • Hugging Face (transformers, accelerate, peft, trl) • NumPy • asyncio • threading • OpenAI • OpenAI Agents SDK • Pydantic • Statsmodels • SciPy • Seaborn • Xgboost • Scikit-learn • Matplotlib • requests • LaTeX • SQL • SLURM

**PUBLICATIONS** | [15+ peer-reviewed \(1,253+ citations\) including 3 first/corresponding-author in top astrophysics journal](#)